

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```


```
In [9]: df = pd.read_csv('sonar_data.csv', header = None)
```

```
In [11]: df
```

```
Out[11]:
```

	0	1	2	3	4	5	6	7	8	9	...	51	52	53	54	55	
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159	0.0072	0.0167	0.
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048	0.0094	0.0191	0.
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095	0.0180	0.0244	0.
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150	0.0085	0.0073	0.
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105	0.0110	0.0015	0.
...
203	0.0187	0.0346	0.0168	0.0177	0.0393	0.1630	0.2028	0.1694	0.2328	0.2684	...	0.0116	0.0098	0.0199	0.0033	0.0101	0.
204	0.0323	0.0101	0.0298	0.0564	0.0760	0.0958	0.0990	0.1018	0.1030	0.2154	...	0.0061	0.0093	0.0135	0.0063	0.0063	0.
205	0.0522	0.0437	0.0180	0.0292	0.0351	0.1171	0.1257	0.1178	0.1258	0.2529	...	0.0160	0.0029	0.0051	0.0062	0.0089	0.
206	0.0303	0.0353	0.0490	0.0608	0.0167	0.1354	0.1465	0.1123	0.1945	0.2354	...	0.0086	0.0046	0.0126	0.0036	0.0035	0.
207	0.0260	0.0363	0.0136	0.0272	0.0214	0.0338	0.0655	0.1400	0.1843	0.2354	...	0.0146	0.0129	0.0047	0.0039	0.0061	0.

208 rows × 61 columns



```
In [13]: df.head()
```

Out[13]:

	0	1	2	3	4	5	6	7	8	9	...	51	52	53	54	55	!
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159	0.0072	0.0167	0.01
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048	0.0094	0.0191	0.01
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095	0.0180	0.0244	0.03
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150	0.0085	0.0073	0.00
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105	0.0110	0.0015	0.00

5 rows × 61 columns

In [15]: `df.describe()`

Out[15]:

	0	1	2	3	4	5	6	7	8	9
count	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000
mean	0.029164	0.038437	0.043832	0.053892	0.075202	0.104570	0.121747	0.134799	0.178003	0.208259
std	0.022991	0.032960	0.038428	0.046528	0.055552	0.059105	0.061788	0.085152	0.118387	0.134416
min	0.001500	0.000600	0.001500	0.005800	0.006700	0.010200	0.003300	0.005500	0.007500	0.011300
25%	0.013350	0.016450	0.018950	0.024375	0.038050	0.067025	0.080900	0.080425	0.097025	0.111275
50%	0.022800	0.030800	0.034300	0.044050	0.062500	0.092150	0.106950	0.112100	0.152250	0.182400
75%	0.035550	0.047950	0.057950	0.064500	0.100275	0.134125	0.154000	0.169600	0.233425	0.268700
max	0.137100	0.233900	0.305900	0.426400	0.401000	0.382300	0.372900	0.459000	0.682800	0.710600

8 rows × 60 columns

In [17]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 208 entries, 0 to 207
Data columns (total 61 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    0      208 non-null    float64
 1    1      208 non-null    float64
 2    2      208 non-null    float64
 3    3      208 non-null    float64
 4    4      208 non-null    float64
 5    5      208 non-null    float64
 6    6      208 non-null    float64
 7    7      208 non-null    float64
 8    8      208 non-null    float64
 9    9      208 non-null    float64
10   10     208 non-null    float64
11   11     208 non-null    float64
12   12     208 non-null    float64
13   13     208 non-null    float64
14   14     208 non-null    float64
15   15     208 non-null    float64
16   16     208 non-null    float64
17   17     208 non-null    float64
18   18     208 non-null    float64
19   19     208 non-null    float64
20   20     208 non-null    float64
21   21     208 non-null    float64
22   22     208 non-null    float64
23   23     208 non-null    float64
24   24     208 non-null    float64
25   25     208 non-null    float64
26   26     208 non-null    float64
27   27     208 non-null    float64
28   28     208 non-null    float64
29   29     208 non-null    float64
30   30     208 non-null    float64
31   31     208 non-null    float64
32   32     208 non-null    float64
33   33     208 non-null    float64
34   34     208 non-null    float64
35   35     208 non-null    float64
36   36     208 non-null    float64

```

```
37 37      208 non-null    float64
38 38      208 non-null    float64
39 39      208 non-null    float64
40 40      208 non-null    float64
41 41      208 non-null    float64
42 42      208 non-null    float64
43 43      208 non-null    float64
44 44      208 non-null    float64
45 45      208 non-null    float64
46 46      208 non-null    float64
47 47      208 non-null    float64
48 48      208 non-null    float64
49 49      208 non-null    float64
50 50      208 non-null    float64
51 51      208 non-null    float64
52 52      208 non-null    float64
53 53      208 non-null    float64
54 54      208 non-null    float64
55 55      208 non-null    float64
56 56      208 non-null    float64
57 57      208 non-null    float64
58 58      208 non-null    float64
59 59      208 non-null    float64
60 60      208 non-null    object
```

dtypes: float64(60), object(1)

memory usage: 99.3+ KB

```
In [19]: df.isnull().sum()
```

```
Out[19]: 0      0
1      0
2      0
3      0
4      0
..
56     0
57     0
58     0
59     0
60     0
```

Length: 61, dtype: int64

```
In [23]: df[60].value_counts()
```


```
Out[23]: 60
M      111
R       97
Name: count, dtype: int64
```

```
In [25]: df.groupby(60).mean()
```

```
Out[25]:
```

	0	1	2	3	4	5	6	7	8	9	...	50	51
60													
M	0.034989	0.045544	0.050720	0.064768	0.086715	0.111864	0.128359	0.149832	0.213492	0.251022	...	0.019352	0.016014
R	0.022498	0.030303	0.035951	0.041447	0.062028	0.096224	0.114180	0.117596	0.137392	0.159325	...	0.012311	0.010453

2 rows × 60 columns



```
In [29]: x = df.drop([60],axis = 1)
```

```
In [31]: y = df[60]
```

```
In [33]: x
```

Out[33]:

	0	1	2	3	4	5	6	7	8	9	...	50	51	52	53	54	
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0232	0.0027	0.0065	0.0159	0.0072	0.
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0125	0.0084	0.0089	0.0048	0.0094	0.
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0033	0.0232	0.0166	0.0095	0.0180	0.
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0241	0.0121	0.0036	0.0150	0.0085	0.
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0156	0.0031	0.0054	0.0105	0.0110	0.
...
203	0.0187	0.0346	0.0168	0.0177	0.0393	0.1630	0.2028	0.1694	0.2328	0.2684	...	0.0203	0.0116	0.0098	0.0199	0.0033	0.
204	0.0323	0.0101	0.0298	0.0564	0.0760	0.0958	0.0990	0.1018	0.1030	0.2154	...	0.0051	0.0061	0.0093	0.0135	0.0063	0.
205	0.0522	0.0437	0.0180	0.0292	0.0351	0.1171	0.1257	0.1178	0.1258	0.2529	...	0.0155	0.0160	0.0029	0.0051	0.0062	0.
206	0.0303	0.0353	0.0490	0.0608	0.0167	0.1354	0.1465	0.1123	0.1945	0.2354	...	0.0042	0.0086	0.0046	0.0126	0.0036	0.
207	0.0260	0.0363	0.0136	0.0272	0.0214	0.0338	0.0655	0.1400	0.1843	0.2354	...	0.0181	0.0146	0.0129	0.0047	0.0039	0.

208 rows × 60 columns



In [35]:

y

Out[35]:

```

0      R
1      R
2      R
3      R
4      R
..
203    M
204    M
205    M
206    M
207    M

```

Name: 60, Length: 208, dtype: object

```
In [37]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size =0.1 ,stratify = y, random_state =1)
```

```
In [39]: a = LogisticRegression()  
a.fit(x_train,y_train)
```

```
Out[39]: ▾ LogisticRegression  
LogisticRegression()
```

```
In [41]: y_train_pred = a.predict(x_train)  
training_data_accuracy = accuracy_score(y_train_pred,y_train)
```

```
In [43]: training_data_accuracy
```

```
Out[43]: 0.8342245989304813
```

```
In [47]: ##accuracy on test data  
y_test_pred = a.predict(x_test)  
test_data_accuracy = accuracy_score(y_test_pred,y_test)
```

```
In [49]: print("Accuracy on test data :", test_data_accuracy)
```

```
Accuracy on test data : 0.7619047619047619
```

```
In [ ]: ##making a predicint system
```

```
In [55]: input_data = (0.0262,0.0582,0.1099,0.1083,0.0974,0.2280,0.2431,0.3771,0.5598,0.6194,0.6333,0.7060,0.5544,0.5320,0.6471)  
  
input_data_as_numpy_array = np.array(input_data)  
#reshape the array  
input_data_reshape = input_data_as_numpy_array.reshape(1,-1)  
  
prediction = a.predict(input_data_reshape)
```

```
In [57]: print("Prediction :", prediction)
```

```
Prediction : ['M']
```

In []: